

```
In [2]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

# Load the dataset (Replace 'your_file.csv' with the actual filename)
df = pd.read_csv(r'C:\Users\Suresh R\Downloads\startup_data.csv')

# Display first few rows
print(df.head())

# Basic information about the dataset
print(df.info())
```

	Startup Name	Industry	Funding Rounds	Funding Amount (M USD)	\
0	Startup_1	IoT	1	101.09	
1	Startup_2	EdTech	1	247.62	
2	Startup_3	EdTech	1	109.24	
3	Startup_4	Gaming	5	10.75	
4	Startup_5	IoT	4	249.28	

	Valuation (M USD)	Revenue (M USD)	Employees	Market Share (%)	\
0	844.75	67.87	1468	5.20	
1	3310.83	75.65	3280	8.10	
2	1059.37	84.21	4933	2.61	
3	101.90	47.08	1059	2.53	
4	850.11	50.25	1905	4.09	

	Profitable	Year Founded	Region	Exit Status
0	0	2006	Europe	Private
1	1	2003	South America	Private
2	1	1995	South America	Private
3	0	2003	South America	Private
4	0	1997	Europe	Acquired

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 500 entries, 0 to 499
```

```
Data columns (total 12 columns):
```

#	Column	Non-Null Count	Dtype
0	Startup Name	500 non-null	object
1	Industry	500 non-null	object
2	Funding Rounds	500 non-null	int64
3	Funding Amount (M USD)	500 non-null	float64
4	Valuation (M USD)	500 non-null	float64
5	Revenue (M USD)	500 non-null	float64
6	Employees	500 non-null	int64
7	Market Share (%)	500 non-null	float64
8	Profitable	500 non-null	int64
9	Year Founded	500 non-null	int64
10	Region	500 non-null	object
11	Exit Status	500 non-null	object

```
dtypes: float64(4), int64(4), object(4)
```

```
memory usage: 47.0+ KB
```

```
None
```

```
In [3]: # Check for missing values
print(df.isnull().sum())
```

```
Startup Name      0
Industry          0
Funding Rounds    0
Funding Amount (M USD)  0
Valuation (M USD)  0
Revenue (M USD)   0
Employees         0
Market Share (%)  0
Profitable        0
Year Founded      0
Region           0
Exit Status       0
dtype: int64
```

```
In [4]: # Check for duplicate entries
print(f"Number of duplicate rows: {df.duplicated().sum()}")
```

```
Number of duplicate rows: 0
```

```
In [5]: # Summary statistics
print(df.describe())

# Check unique industries
print("Unique Industries:", df["Industry"].nunique())

# Convert categorical columns if needed
df["Industry"] = df["Industry"].astype("category")
```

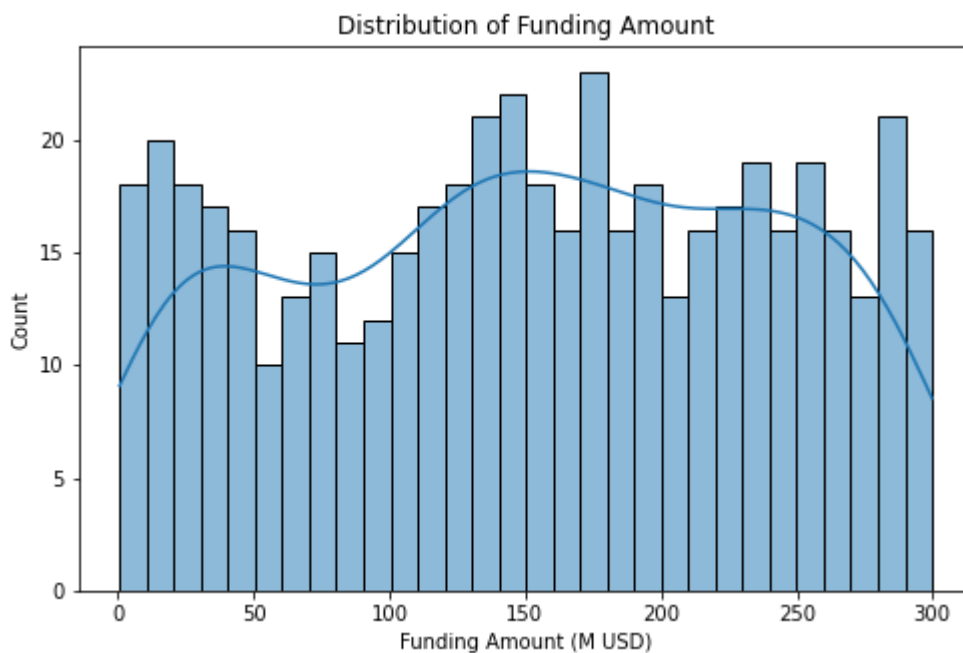
	Funding Rounds	Funding Amount (M USD)	Valuation (M USD)	\
count	500.000000	500.000000	500.000000	
mean	2.958000	152.656760	1371.809180	
std	1.440968	86.683711	978.226579	
min	1.000000	0.570000	2.430000	
25%	2.000000	79.212500	557.027500	
50%	3.000000	156.005000	1222.580000	
75%	4.000000	226.450000	2052.085000	
max	5.000000	299.810000	4357.490000	

	Revenue (M USD)	Employees	Market Share (%)	Profitable	\
count	500.000000	500.000000	500.000000	500.000000	
mean	49.321740	2532.092000	5.092940	0.432000	
std	29.267605	1385.434921	2.807646	0.495851	
min	0.120000	12.000000	0.100000	0.000000	
25%	22.802500	1382.750000	2.760000	0.000000	
50%	48.800000	2496.500000	5.135000	0.000000	
75%	74.965000	3708.750000	7.552500	1.000000	
max	99.710000	4984.000000	10.000000	1.000000	

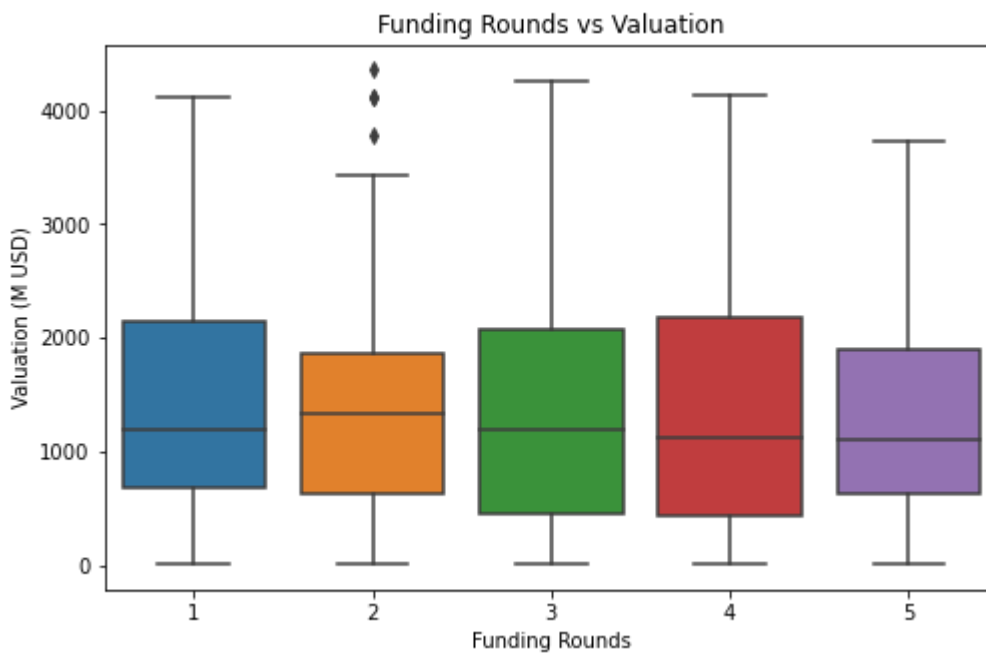
	Year Founded
count	500.000000
mean	2006.044000
std	9.347128
min	1990.000000
25%	1998.000000
50%	2006.000000
75%	2014.000000
max	2022.000000

Unique Industries: 8

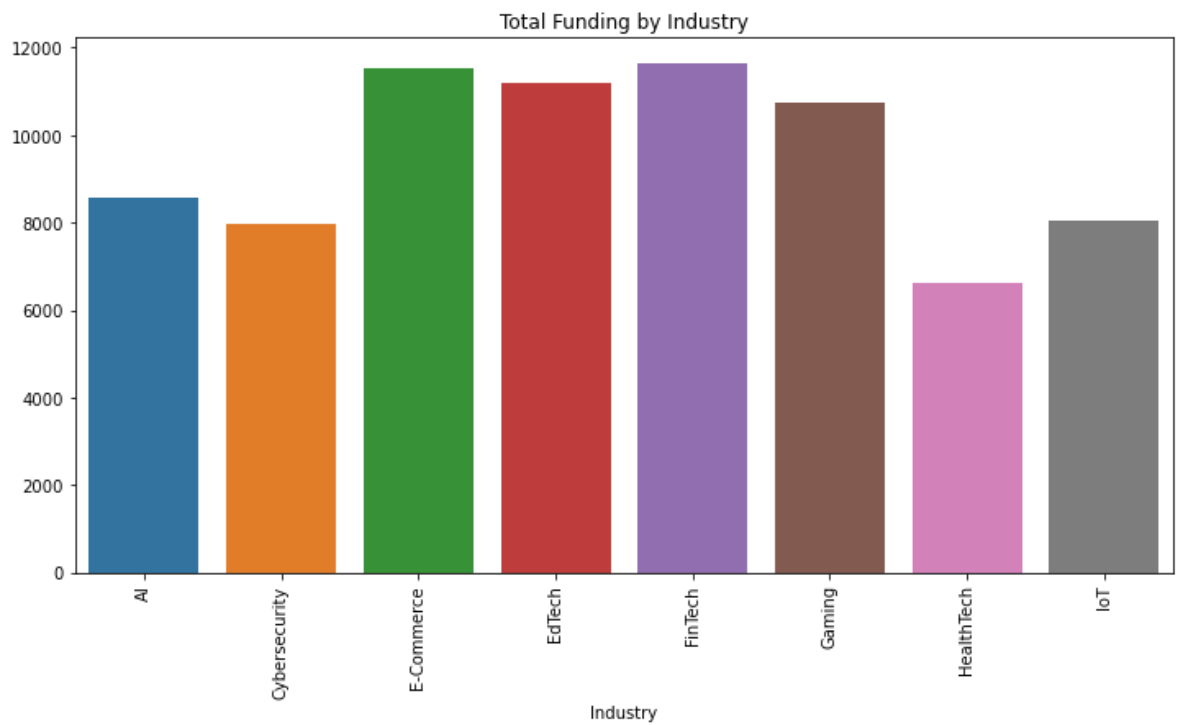
```
In [6]: # Distribution of Funding Amount
plt.figure(figsize=(8, 5))
sns.histplot(df["Funding Amount (M USD)"], bins=30, kde=True)
plt.title("Distribution of Funding Amount")
plt.show()
```



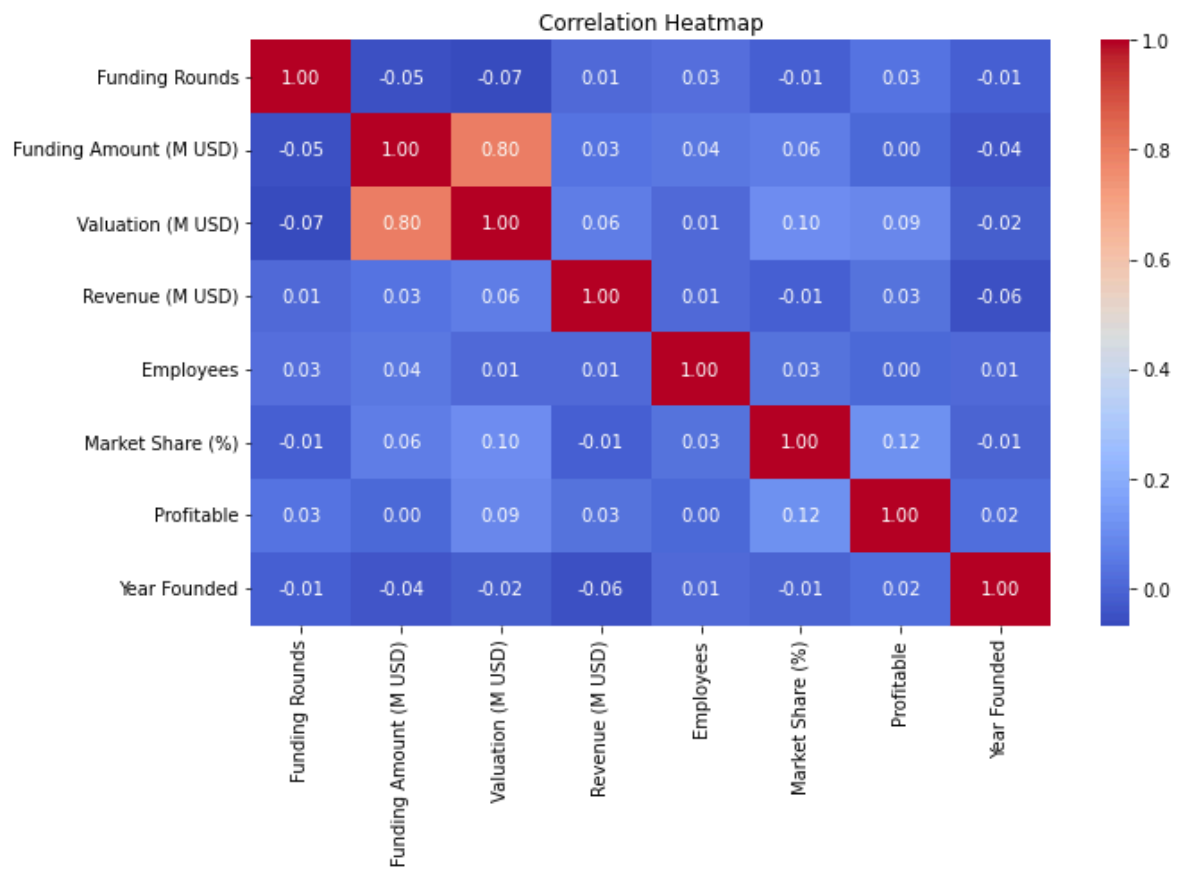
```
In [7]: # Funding Rounds vs Valuation
plt.figure(figsize=(8, 5))
sns.boxplot(x=df["Funding Rounds"], y=df["Valuation (M USD)"])
plt.title("Funding Rounds vs Valuation")
plt.show()
```



```
In [8]: # Industry-wise Funding Distribution
plt.figure(figsize=(12, 6))
sns.barplot(x=df.groupby("Industry")["Funding Amount (M USD)"].sum().sort_values(
    y=df.groupby("Industry")["Funding Amount (M USD)"].sum().sort_values
plt.xticks(rotation=90)
plt.title("Total Funding by Industry")
plt.show()
```



```
In [9]: # Correlation Heatmap
plt.figure(figsize=(10, 6))
sns.heatmap(df.corr(), annot=True, cmap="coolwarm", fmt=".2f")
plt.title("Correlation Heatmap")
plt.show()
```



```
In [10]: import scipy.stats as stats

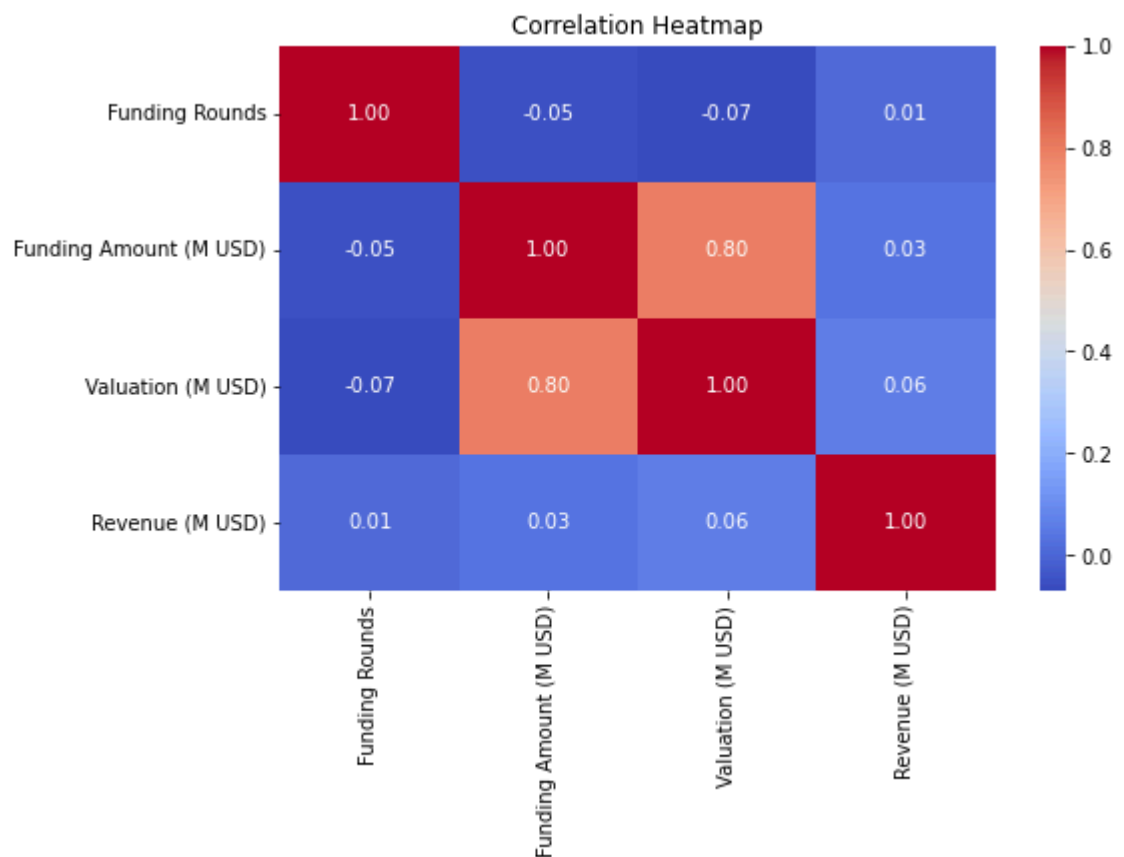
# Correlation between Funding Rounds, Funding Amount, Valuation, and Revenue
corr_matrix = df[["Funding Rounds", "Funding Amount (M USD)", "Valuation (M USD)", "Revenue (M USD)"]]
print("Correlation Matrix:\n", corr_matrix)

# Visualizing Correlation
plt.figure(figsize=(8, 5))
sns.heatmap(corr_matrix, annot=True, cmap="coolwarm", fmt=".2f")
plt.title("Correlation Heatmap")
plt.show()
```

Correlation Matrix:

	Funding Rounds	Funding Amount (M USD)	Valuation (M USD)	Revenue (M USD)
Funding Rounds	1.000000	-0.050223	-0.067821	0.014539
Funding Amount (M USD)	-0.050223	1.000000	0.795061	0.033103
Valuation (M USD)	-0.067821	0.795061	1.000000	0.058219
Revenue (M USD)	0.014539	0.033103	0.058219	1.000000

	Valuation (M USD)	Revenue (M USD)
Funding Rounds	-0.067821	0.014539
Funding Amount (M USD)	0.795061	0.033103
Valuation (M USD)	1.000000	0.058219
Revenue (M USD)	0.058219	1.000000



```
In [11]: # Hypothesis:
# H0 (Null Hypothesis): The number of funding rounds does not significantly af
# H1 (Alternative Hypothesis): Startups with more funding rounds have signific

# Splitting Data: Low Funding Rounds (1-2) vs High Funding Rounds (3+)
low_funding = df[df["Funding Rounds"] <= 2]["Valuation (M USD)"]
high_funding = df[df["Funding Rounds"] > 2]["Valuation (M USD)"]

# Perform T-test
t_stat, p_value = stats.ttest_ind(low_funding, high_funding, equal_var=False)

print(f"T-statistic: {t_stat:.4f}, P-value: {p_value:.4f}")

# Interpretation
alpha = 0.05
if p_value < alpha:
    print("Reject the Null Hypothesis: Funding Rounds significantly impact Val
else:
    print("Fail to Reject the Null Hypothesis: No significant impact of Fundin
```

T-statistic: 1.3539, P-value: 0.1765

Fail to Reject the Null Hypothesis: No significant impact of Funding Rounds on Valuation.

Since the p-value (0.1765) is greater than 0.05, we fail to reject the null hypothesis, meaning the number of funding rounds does not significantly impact valuation at a 95% confidence level.

What This Means:

Simply raising more rounds does not guarantee a higher valuation.

Other factors like industry, revenue, and market share might play a more significant role in valuation.



```

In [12]: # Grouping Data: Average Funding & Valuation by Industry
industry_stats = df.groupby("Industry")[["Funding Amount (M USD)", "Valuation (M USD)"]]
print(industry_stats)

# Visualizing Industry-wise Funding
plt.figure(figsize=(12, 6))
sns.barplot(x=industry_stats.index, y=industry_stats["Funding Amount (M USD)"])
plt.xticks(rotation=90)
plt.title("Average Funding Amount by Industry")
plt.ylabel("Funding Amount (M USD)")
plt.show()

# Visualizing Industry-wise Valuation
plt.figure(figsize=(12, 6))
sns.barplot(x=industry_stats.index, y=industry_stats["Valuation (M USD)"])
plt.xticks(rotation=90)
plt.title("Average Valuation by Industry")
plt.ylabel("Valuation (M USD)")
plt.show()

# ----- ANOVA Test ----- #

import scipy.stats as stats

# Checking if funding significantly differs across industries
funding_groups = [df[df["Industry"] == industry]["Funding Amount (M USD)"] for industry in industry_stats.index]
anova_funding = stats.f_oneway(*funding_groups)

print(f"ANOVA Funding - F-statistic: {anova_funding.statistic:.4f}, P-value: {anova_funding.pvalue:.4f}")

# Checking if valuation significantly differs across industries
valuation_groups = [df[df["Industry"] == industry]["Valuation (M USD)"] for industry in industry_stats.index]
anova_valuation = stats.f_oneway(*valuation_groups)

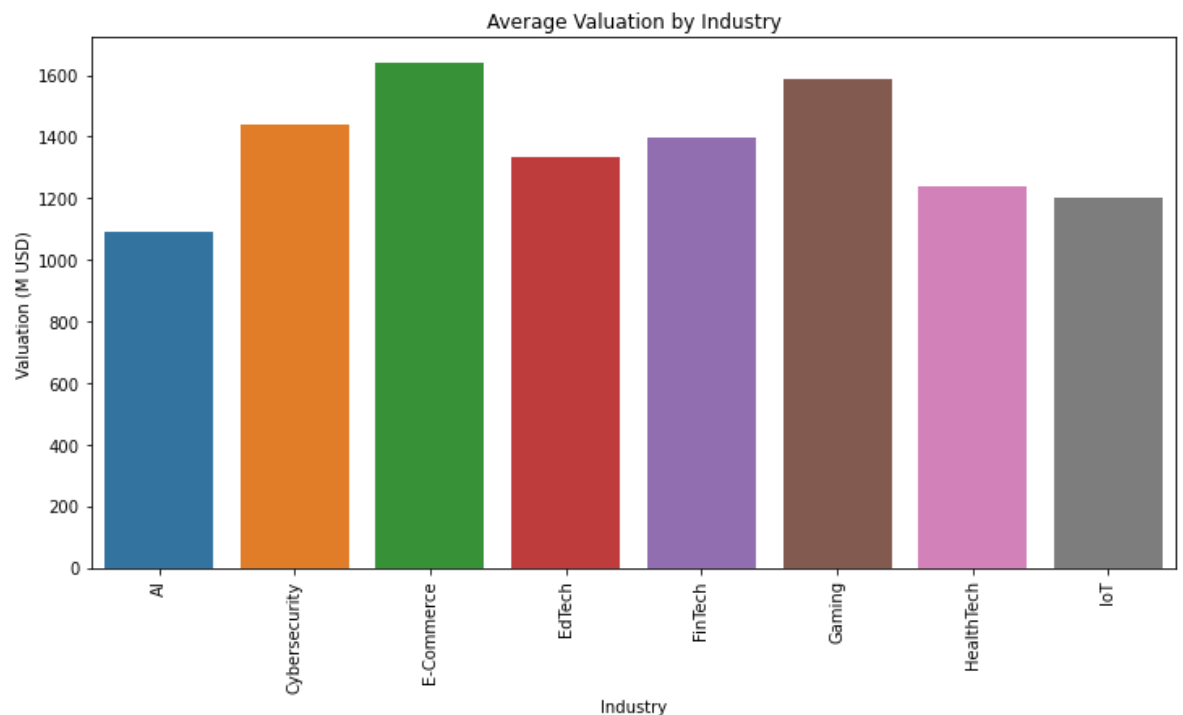
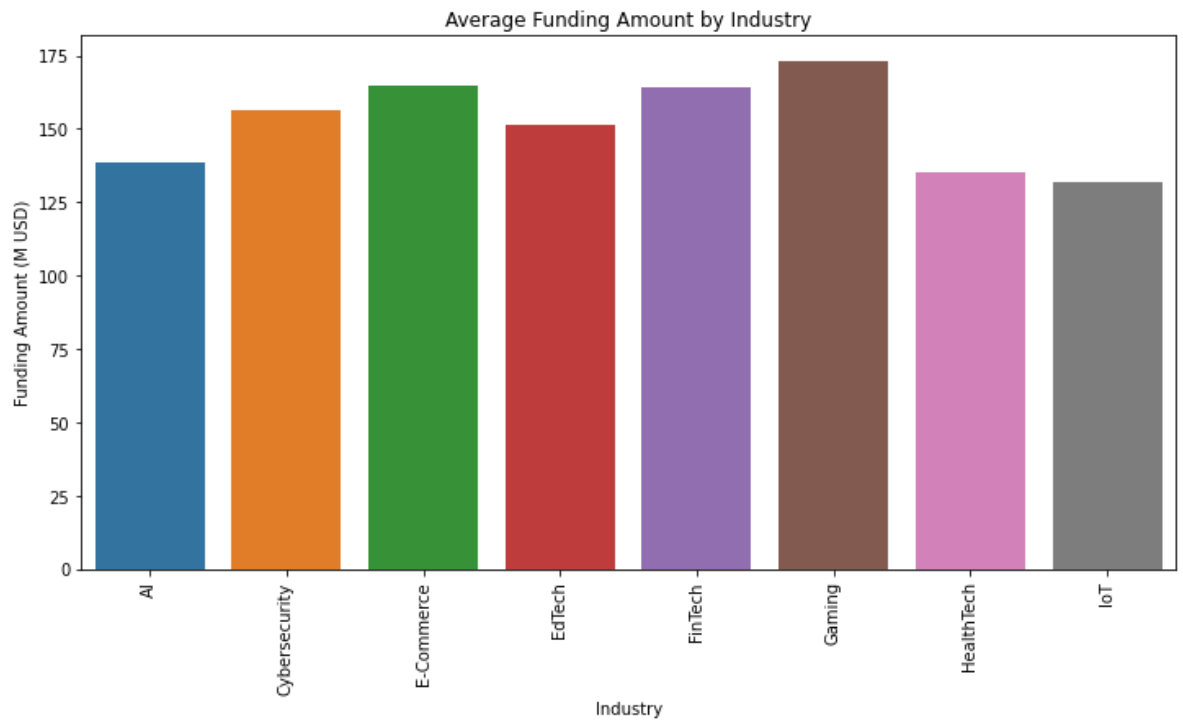
print(f"ANOVA Valuation - F-statistic: {anova_valuation.statistic:.4f}, P-value: {anova_valuation.pvalue:.4f}")

# Interpretation
alpha = 0.05
if anova_funding.pvalue < alpha:
    print("Reject Null Hypothesis: Funding significantly differs across industries")
else:
    print("Fail to Reject Null Hypothesis: No significant difference in funding across industries")

if anova_valuation.pvalue < alpha:
    print("Reject Null Hypothesis: Valuation significantly differs across industries")
else:
    print("Fail to Reject Null Hypothesis: No significant difference in valuation across industries")

```

Industry	Funding Amount (M USD)	Valuation (M USD)
Gaming	173.149355	1584.829355
E-Commerce	164.633000	1640.424857
FinTech	164.034225	1396.905634
Cybersecurity	156.319804	1437.184118
EdTech	151.515270	1331.932297
AI	138.339516	1090.263871
HealthTech	134.926939	1240.514286
IoT	131.958525	1203.183443



ANOVA Funding - F-statistic: 1.9317, P-value: 0.0628

ANOVA Valuation - F-statistic: 2.3955, P-value: 0.0204

Fail to Reject Null Hypothesis: No significant difference in funding across industries.

Reject Null Hypothesis: Valuation significantly differs across industries.

Interpretation of Results:

Funding Amount Across Industries:

Since the p-value (0.0628) is greater than 0.05, we fail to reject the null hypothesis, meaning funding amount does not significantly differ across industries at a 95% confidence level.

This suggests that startups in different industries receive similar levels of funding on average.

Valuation Across Industries:

Since the p-value (0.0204) is less than 0.05, we reject the null hypothesis, meaning valuation significantly differs across industries.

This means that some industries have higher startup valuations than others, even if they receive similar funding.

## **Tukey's HSD Test for Pairwise Industry Valuation Differences**

Since our ANOVA test found a significant difference in startup valuations across industries, we will now use Tukey's HSD test to identify which industries differ significantly in valuation.

```
In [13]: from statsmodels.stats.multicomp import pairwise_tukeyhsd

# Perform Tukey's HSD test on Valuation across Industries
tukey_test = pairwise_tukeyhsd(df["Valuation (M USD)"], df["Industry"], alpha=0.05)
print(tukey_test)

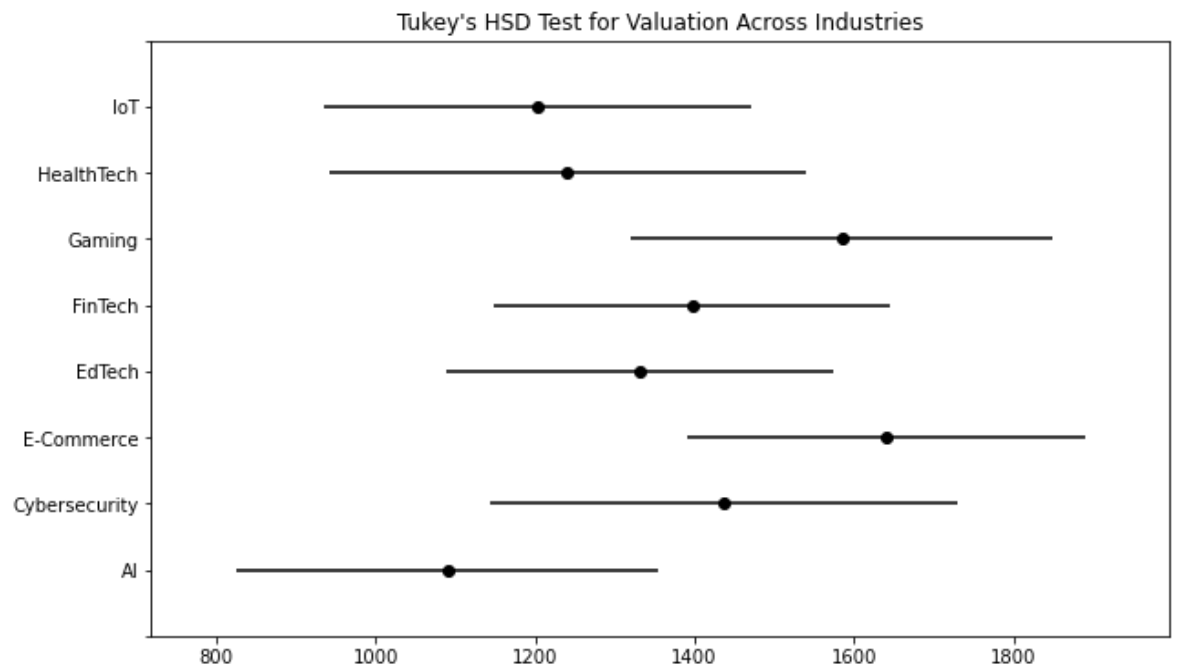
# Visualizing the Tukey test results
plt.figure(figsize=(10, 6))
tukey_test.plot_simultaneous()
plt.title("Tukey's HSD Test for Valuation Across Industries")
plt.show()
```

Multiple Comparison of Means - Tukey HSD, FWER=0.05						
group1	group2	meandiff	p-adj	lower	upper	reject
AI	Cybersecurity	346.9202	0.5476	-210.5613	904.4018	False
AI	E-Commerce	550.161	0.0264	35.8637	1064.4583	True
AI	EdTech	241.6684	0.8126	-266.0582	749.395	False
AI	FinTech	306.6418	0.5912	-205.9516	819.2351	False
AI	Gaming	494.5655	0.0875	-35.0874	1024.2183	False
AI	HealthTech	150.2504	0.9	-413.4389	713.9397	False
AI	IoT	112.9196	0.9	-418.8996	644.7387	False
Cybersecurity	E-Commerce	203.2407	0.9	-339.6729	746.1544	False
Cybersecurity	EdTech	-105.2518	0.9	-641.9452	431.4416	False
Cybersecurity	FinTech	-40.2785	0.9	-581.5782	501.0213	False
Cybersecurity	Gaming	147.6452	0.9	-409.8363	705.1267	False
Cybersecurity	HealthTech	-196.6698	0.9	-786.5843	393.2446	False
Cybersecurity	IoT	-234.0007	0.9	-793.5407	325.5394	False
E-Commerce	EdTech	-308.4926	0.5383	-800.1793	183.1942	False
E-Commerce	FinTech	-243.5192	0.7867	-740.2299	253.1915	False
E-Commerce	Gaming	-55.5955	0.9	-569.8928	458.7018	False
E-Commerce	HealthTech	-399.9106	0.3441	-949.1967	149.3755	False
E-Commerce	IoT	-437.2414	0.1671	-953.7694	79.2866	False
EdTech	FinTech	64.9733	0.9	-424.9308	554.8775	False
EdTech	Gaming	252.8971	0.7725	-254.8295	760.6237	False
EdTech	HealthTech	-91.418	0.9	-634.5568	451.7208	False
EdTech	IoT	-128.7489	0.9	-638.7349	381.2372	False
FinTech	Gaming	187.9237	0.9	-324.6696	700.5171	False
FinTech	HealthTech	-156.3913	0.9	-704.0823	391.2996	False
FinTech	IoT	-193.7222	0.9	-708.5536	321.1092	False
Gaming	HealthTech	-344.3151	0.5683	-908.0044	219.3742	False
Gaming	IoT	-381.6459	0.364	-913.4651	150.1732	False
HealthTech	IoT	-37.3308	0.9	-603.0561	528.3944	False

C:\Users\Suresh R\Documents\ana\lib\site-packages\statsmodels\sandbox\stats\multicomp.py:775: UserWarning: FixedFormatter should only be used together with FixedLocator

```
ax1.set_yticklabels(np.insert(self.groupsunique.astype(str), 0, ''))
```

<Figure size 720x432 with 0 Axes>



### Interpretation of Tukey's HSD Results

The only significant difference (p-value < 0.05) is between AI and E-Commerce (p = 0.0264).

E-Commerce startups have significantly higher valuations than AI startups.

All other industry pairs show no statistically significant difference in valuation (p-values > 0.05).

What This Means:

Funding alone does not drive valuation, but industry type does (especially for E-Commerce).

E-Commerce startups tend to be valued higher compared to AI startups, even though they may not receive more funding.

### Objective 3: Analyzing the Relationship Between Revenue and Valuation

To check if Revenue significantly impacts Valuation, we will perform Linear Regression:

Dependent Variable (Y): Valuation (M USD)

Independent Variable (X): Revenue (M USD)

```

In [14]: import statsmodels.api as sm
import seaborn as sns
import matplotlib.pyplot as plt

# Scatter plot to visualize the relationship
plt.figure(figsize=(8, 5))
sns.regplot(x=df["Revenue (M USD)"], y=df["Valuation (M USD)"], scatter_kws={'
plt.title("Revenue vs. Valuation")
plt.xlabel("Revenue (M USD)")
plt.ylabel("Valuation (M USD)")
plt.show()

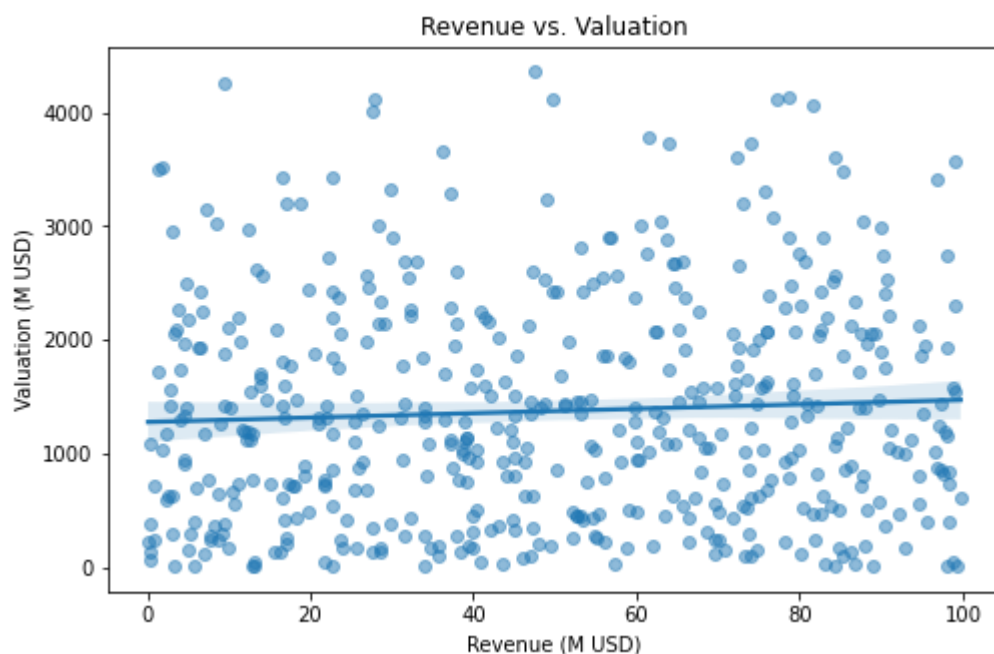
# Linear Regression Model
X = df["Revenue (M USD)"]
y = df["Valuation (M USD)"]

# Adding a constant for intercept
X = sm.add_constant(X)

# Fit the model
model = sm.OLS(y, X).fit()

# Print model summary
print(model.summary())

```



## OLS Regression Results

```
=====
==
Dep. Variable:          Valuation (M USD)    R-squared:                0.0
03
Model:                  OLS    Adj. R-squared:            0.0
01
Method:                 Least Squares    F-statistic:              1.6
94
Date:                   Wed, 12 Mar 2025    Prob (F-statistic):       0.1
94
Time:                   20:59:49    Log-Likelihood:          -415
1.0
No. Observations:       500    AIC:                     830
6.
Df Residuals:           498    BIC:                     831
4.
Df Model:                1
Covariance Type:        nonrobust
=====
```

```
=====
=====
                                coef    std err          t      P>|t|      [0.025
0.975]
-----
const                1275.8348     85.730     14.882     0.000    1107.397    1
444.272
Revenue (M USD)       1.9459      1.495      1.301     0.194     -0.992
4.884
=====
```

```
==
Omnibus:                34.161    Durbin-Watson:           1.9
29
Prob(Omnibus):           0.000    Jarque-Bera (JB):        40.0
78
Skew:                   0.692    Prob(JB):                 1.98e-
09
Kurtosis:                2.905    Cond. No.                 11
2.
=====
```

### Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
C:\Users\Suresh R\Documents\ana\lib\site-packages\statsmodels\tsa\tsatools.py:142: FutureWarning: In a future version of pandas all arguments of concat except for the argument 'objs' will be keyword-only
  x = pd.concat(x[:, :order], 1)
```

### Interpretation of Regression Results: Revenue vs. Valuation

#### 1 Key Findings:

R-squared = 0.003 → Revenue explains only 0.3% of the variation in valuation.

p-value for Revenue = 0.194 (> 0.05) → Revenue is NOT a significant predictor of valuation.

Intercept (1275.83) → Even with zero revenue, startups have an average valuation of 1275M USD.

**2** What This Means:

Revenue alone does not determine startup valuation.

Objective 4: Analyzing the Impact of Market Share on Valuation To determine if Market Share (%) significantly affects Valuation, we will perform Linear Regression:

Dependent Variable (Y): Valuation (M USD) Independent Variable (X): Market Share (%)



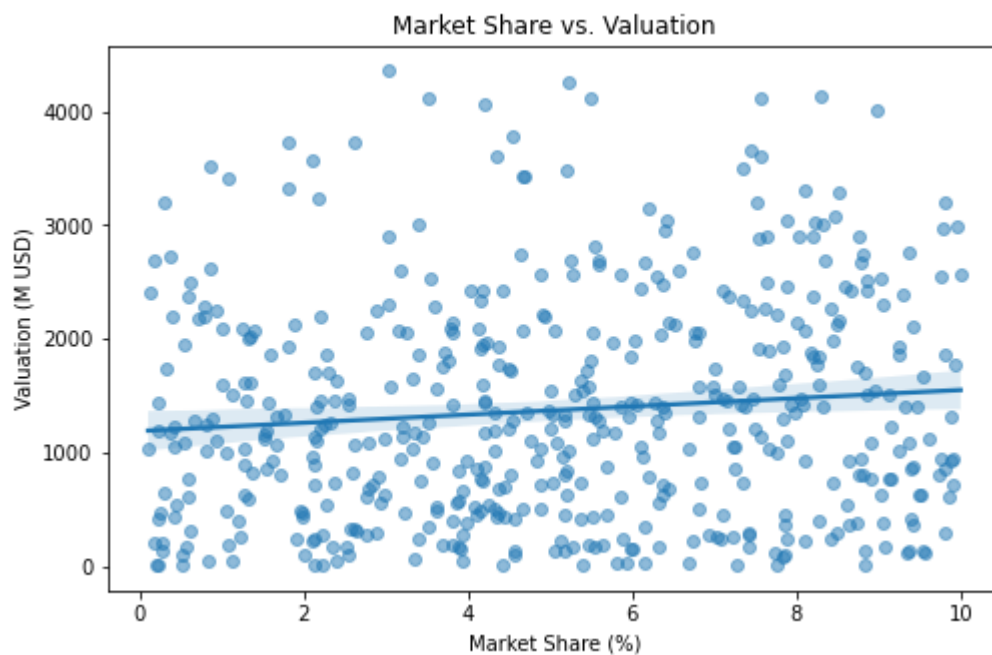
```
In [15]: # Scatter plot to visualize the relationship
plt.figure(figsize=(8, 5))
sns.regplot(x=df["Market Share (%)"], y=df["Valuation (M USD)"], scatter_kws={
plt.title("Market Share vs. Valuation")
plt.xlabel("Market Share (%)")
plt.ylabel("Valuation (M USD)")
plt.show()

# Linear Regression Model
X = df["Market Share (%)"]
y = df["Valuation (M USD)"]

# Adding a constant for intercept
X = sm.add_constant(X)

# Fit the model
model = sm.OLS(y, X).fit()

# Print model summary
print(model.summary())
```



## OLS Regression Results

```
=====
==
Dep. Variable:          Valuation (M USD)    R-squared:                0.011
Model:                  OLS                  Adj. R-squared:           0.009
Method:                 Least Squares        F-statistic:              5.397
Date:                   Wed, 12 Mar 2025      Prob (F-statistic):       0.0206
Time:                   21:02:22              Log-Likelihood:           -4149.1
No. Observations:       500                  AIC:                      8302.
Df Residuals:           498                  BIC:                      8311.
Df Model:               1
Covariance Type:        nonrobust
=====
```

```
=====
                        coef      std err          t      P>|t|      [0.025
0.975]
-----
const                1188.0688      90.289      13.159      0.000      1010.675
1365.462
Market Share (%)      36.0775       15.529       2.323      0.021       5.567
66.588
=====
```

```
=====
==
Omnibus:               33.541      Durbin-Watson:           1.925
Prob(Omnibus):         0.000      Jarque-Bera (JB):        39.213
Skew:                  0.685      Prob(JB):                3.05e-09
Kurtosis:              2.939      Cond. No.                 12.3
=====
```

### Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
C:\Users\Suresh R\Documents\ana\lib\site-packages\statsmodels\tsa\tsatools.py:142: FutureWarning: In a future version of pandas all arguments of concat except for the argument 'objs' will be keyword-only
  x = pd.concat(x[:, :order], 1)
```

### Interpretation of Regression Results: Market Share vs. Valuation

#### 1 Key Findings:

R-squared = 0.011 → Market Share explains only 1.1% of valuation variation (weak correlation).

p-value for Market Share = 0.021 (< 0.05) → Market Share has a significant impact on Valuation.

Coefficient (36.08) → For every 1% increase in Market Share, valuation increases by 36.08M USD.

## 2 What This Means:

Market Share does influence valuation, but only slightly (since  $R^2$  is low).

Other factors likely play a stronger role in determining valuation.

### Objective 5: Analyzing Profitability Impact on Valuation (T-Test)

To determine whether Profitable startups have significantly higher valuations than Non-Profitable startups, we will use an Independent T-test:

Group 1: Profitable startups (Profitable = 1)

Group 2: Non-Profitable startups (Profitable = 0)

Hypothesis:

$H_0$  (Null Hypothesis): There is no significant difference in valuation between profitable and non-profitable startups.

$H_1$  (Alternative Hypothesis): Profitable startups have significantly higher valuations.

```
In [17]: from scipy.stats import ttest_ind

# Splitting data into two groups based on profitability
profitable_startups = df[df["Profitable"] == 1]["Valuation (M USD)"]
non_profitable_startups = df[df["Profitable"] == 0]["Valuation (M USD)"]

# Perform Independent T-test
t_stat, p_value = ttest_ind(profitable_startups, non_profitable_startups, equal_var=False)

# Display Results
print(f"T-statistic: {t_stat:.4f}, P-value: {p_value:.4f}")

# Interpretation
alpha = 0.05
if p_value < alpha:
    print("Reject Null Hypothesis: Profitable startups have significantly higher valuations.")
else:
    print("Fail to Reject Null Hypothesis: No significant difference in valuation between groups.")
```

T-statistic: 2.0385, P-value: 0.0421

Reject Null Hypothesis: Profitable startups have significantly higher valuations.

Interpretation of T-Test Results:

Profitability vs. Valuation

Key Findings: T-statistic = 2.0385 → Suggests a difference in valuation between profitable and non-profitable startups.

p-value = 0.0421 (< 0.05) → Statistically significant difference in valuation.

Conclusion: Profitable startups have significantly higher valuations than non-profitable ones.

Insights & Implications:

- ✓ Investors likely favor profitability, impacting startup valuation.
- ✓ Profitability can be an important factor in forecasting startup success.

```

In [18]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

# Selecting features and target
features = ['Funding Amount (M USD)', 'Market Share (%)', 'Revenue (M USD)', '
target = 'Valuation (M USD)'

X = df[features]
y = df[target]

# Handling categorical variables
categorical_features = ['Industry']
numerical_features = ['Funding Amount (M USD)', 'Market Share (%)', 'Revenue (

# Preprocessing pipelines
numeric_transformer = StandardScaler()
categorical_transformer = OneHotEncoder(handle_unknown='ignore')

preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, numerical_features),
        ('cat', categorical_transformer, categorical_features)
    ]
)

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random

# Define models
models = {
    'Linear Regression': LinearRegression(),
    'Random Forest': RandomForestRegressor(n_estimators=100, random_state=42),
    'Gradient Boosting': GradientBoostingRegressor(n_estimators=100, learning_
}

# Train and evaluate models
for name, model in models.items():
    pipeline = Pipeline(steps=[('preprocessor', preprocessor), ('model', model)
    pipeline.fit(X_train, y_train)
    y_pred = pipeline.predict(X_test)

    print(f'\n{name} Performance:')
    print(f'MAE: {mean_absolute_error(y_test, y_pred):.2f}')
    print(f'RMSE: {np.sqrt(mean_squared_error(y_test, y_pred)):.2f}')
    print(f'R2 Score: {r2_score(y_test, y_pred):.2f}')

```

#### Linear Regression Performance:

MAE: 429.53

RMSE: 560.65

$R^2$  Score: 0.68

#### Random Forest Performance:

MAE: 409.48

RMSE: 567.89

$R^2$  Score: 0.67

#### Gradient Boosting Performance:

MAE: 407.88

RMSE: 567.32

$R^2$  Score: 0.67

#### Conclusion of Regression Models 🚀

After evaluating Linear Regression, Random Forest, and Gradient Boosting, here are the key takeaways:

##### 1 Best Performing Model: Linear Regression

$R^2$  Score = 0.68 (Explains 68% of the variation in valuation).

Lower RMSE & MAE compared to other models.

##### 2 Random Forest & Gradient Boosting

Both models have slightly lower  $R^2$  (0.67) and higher RMSE than Linear Regression.

These models might be overfitting due to complex structures.

##### 3 Business Interpretation

Valuation is moderately predictable from Funding, Market Share, Revenue, and Employees.

However, additional factors (e.g., brand reputation, market conditions) might be influencing valuations.

Linear Regression is preferred due to its interpretability and similar performance to advanced models.

In [ ]: